

THE MIDNIGHT ARCHITECT

Complete Job Search & Personal Branding Toolkit

12 ready-to-use career assets — built in one Claude session from a single resume upload. Cover letters, outreach emails, brand framework, skill gap analysis, target company strategy, and interview preparation.

Deborah Olofin

Data Scientist & Machine Learning Engineer

github.com/Mi-kami | linkedin.com/in/deborah-olofin | deboraholofin43@gmail.com

12 Sections	5 Target Companies	10 Job Titles	2 Live ML Systems
-------------	--------------------	---------------	-------------------

To: hello@indexzoo.com | Subject: *Deborah Olofin - Olabisi Onabanjo University - Nigeria - ZF Intern*

Dear Zoo Finance Hiring Team,

I am writing to apply for the Data Scientist (AI Focus) Intern position at Zoo Finance. My background maps directly onto what this role requires: a production-deployed Crypto AI Intelligence System, demonstrated experience processing and modeling real-world financial time-series data, and an AI-augmented workflow built around the exact tools your JD references — Claude, Python, and AI-assisted code development.

My most relevant project is the Crypto AI Intelligence System I built and deployed independently. Phases 1 and 2 — live multi-source cryptocurrency data ingestion and pipeline architecture — are deployed and running on Railway with Docker, Streamlit, SQLite, and CI/CD. Phase 3 (ML prediction pipeline) is in active development. This is not a tutorial project — it is a production system I engineered from scratch.

Before that, I served as Team Lead and Lead Data Scientist on the Early Warning Food Price Monitor for Northeast Nigeria — a project that required engineering a clean master dataset from over 50,000 raw rows across 13 states, 7 commodities, and 8 years, integrating sources from WFP, CBN, NBS, ACLED, NIMET, and FAO. I built and evaluated SARIMA, SARIMAX, XGBoost, and Facebook Prophet models, resulting in a live Streamlit deployment with 11 of 13 commodity-state combinations production-ready and a best-case MAPE of 14.7%.

I use AI tools as genuine productivity multipliers in my daily workflow. Claude, GitHub Copilot, and Jupyter AI are active parts of how I debug, prototype, and generate analysis — not novelties. I am also deeply comfortable in async, fully remote environments, having completed my Codveda Technologies internship entirely remotely, delivering six production-quality ML projects.

Zoo Finance sits at an intersection I find analytically compelling: rigorous data science applied to decentralized finance. I am ready to contribute to Web3 market data analysis, predictive modeling for DeFi protocols, and competitor intelligence from day one, and I am actively deepening my on-chain tooling knowledge alongside this application.

I would welcome the opportunity to speak with your team.

Warm regards,

Deborah Olofin

deboraholofin43@gmail.com | github.com/Mi-kami | linkedin.com/in/deborah-olofin

Subject: *Data Scientist Intern (AI Focus) — CS Graduate | Live Crypto ML System | Python · Docker · AI-Augmented Workflow*

Hi [Recruiter Name],

I came across the Data Scientist (AI Focus) Intern role at Zoo Finance and wanted to reach out directly before submitting my application. A few things that may be relevant to you:

- I independently built and deployed a live Crypto AI Intelligence System — multi-source crypto data ingestion pipeline (Phases 1–2 live on Railway), ML prediction pipeline in active development. Stack: Python, SQLite, MLflow, Docker, Streamlit, CI/CD.
- I led the data science work on a production food price forecasting system — 50,000+ raw rows across 13 states, 7 commodities, 8 years, live Streamlit deployment, MAPE as low as 14.7%.
- My daily workflow includes Claude, GitHub Copilot, and Jupyter AI — exactly the AI-assisted toolkit your JD references.
- I hold a BSc in Computer Science (3.78 CGPA) and am deepening ML skills through DataCamp, AWS AI & ML Scholars, and WorldQuant's Applied Data Science Lab.

I am genuinely interested in the Web3 and DeFi data space — not as a buzzword but because it is what I am actively building in. Would you be open to a 15-minute conversation?

Deborah Olofin

deboraholofin43@gmail.com | github.com/Mi-kami | linkedin.com/in/deborah-olofin

03

HIRING MANAGER EMAIL

Subject: Data Scientist Intern Application — I'm Already Building What You're Hiring For

Dear [Hiring Manager's Name],

I'm applying for the Data Scientist (AI Focus) Intern position at Zoo Finance, and I wanted to write to you directly because the work your team is doing is not abstract to me — it is what I am actively building.

My Crypto AI Intelligence System is a live, independent project: multi-source cryptocurrency data ingestion deployed on Railway with Docker and CI/CD (Phases 1–2), with the ML prediction pipeline in Phase 3 development. I designed the full architecture — from raw data ingestion through storage (SQLite), experiment tracking (MLflow), and visualization (Streamlit). This is not a bootcamp demo. It is a production system I own entirely.

Prior to that, I led data science on a food price forecasting project from 50,000+ raw rows across 13 states, 7 commodities, and 8 data sources to a live Streamlit deployment — with SARIMA, SARIMAX, XGBoost, and Prophet evaluated head-to-head. Best-case MAPE: 14.7%.

I work AI-augmented natively. My CV and GitHub are linked below. I would welcome the chance to discuss how I can contribute directly to Zoo Finance's analytics work.

Deborah Olofin

deboraholofin43@gmail.com | github.com/Mi-kami

04

LINKEDIN CONNECTION REQUEST (MAX 300 CHARACTERS)

Hi [Name] — CS graduate with a live Crypto AI pipeline (deployed on Railway) and production forecasting experience. Excited about Zoo Finance's DeFi data work. My background sits right at this intersection. Would love to connect.

05

REFERRAL REQUEST MESSAGE

Subject: Quick Favour — Zoo Finance Intern Referral

Hi [Name], Hope you're doing well! I came across the Data Scientist (AI Focus) Intern role at Zoo Finance and I'm genuinely excited about it — the Web3 data science work they're doing is a direct match for what I've been building independently.

Quick context: Live Crypto AI Intelligence System (multi-source crypto data pipeline, production on Railway) + led food price forecasting — 50K+ rows, 6 sources, live deployment, 14.7% best MAPE. Stack: Python, MLflow, Docker, Streamlit, CI/CD; AI tools: Claude, GitHub Copilot daily.

If you have any connection to the Zoo Finance team, an introduction would mean a lot. Happy to share my GitHub and resume first so you can decide whether it makes sense. Either way, thank you.

Deborah Olofin

github.com/Mi-kami | linkedin.com/in/deborah-olofin

06

FOLLOW-UP EMAIL (DAY 5 — NO RESPONSE)

Subject: Following Up — Data Scientist Intern Application | Deborah Olofin

Dear [Name / Hiring Team], I'm following up on my application for the Data Scientist (AI Focus) Intern role at Zoo Finance, submitted [X days] ago.

My interest hasn't changed. My Crypto AI Intelligence System — a live multi-source crypto data pipeline with ML integration in active development — is directly relevant to the on-chain/off-chain analysis your team is focused on, and I would welcome the chance to demonstrate that in a conversation.

If helpful, I'm happy to share a walkthrough of the GitHub repo or answer technical questions upfront. I appreciate your time either way.

Deborah Olofin

deboraholofin43@gmail.com | github.com/Mi-kami

Use at: Interviews | Networking events | Discord/Twitter Spaces | Career fairs

"I'm Deborah Olofin — a data scientist and machine learning engineer with a Computer Science background and two production ML systems already in deployment.

My most recent independent work is a live Crypto AI Intelligence System — a multi-source cryptocurrency data ingestion and intelligence pipeline running in production on Railway, with ML prediction integration in active development. Before that, I served as team lead and lead data scientist on a food price forecasting project that processed 50,000-plus raw rows from 6 real-world data sources across 13 states and 8 years — and went end-to-end into a live Streamlit deployment with a best-case forecast error of 14.7%.

I work across the full ML stack — data engineering, feature engineering, model training, evaluation, and deployment — in Python, with MLflow, Docker, Streamlit, and CI/CD. AI tools — Claude, GitHub Copilot — are native to my workflow.

I'm currently focused on AI and ML roles in data-intensive, high-impact environments, with particular interest in the Web3 and DeFi analytics space."

#	Job Title	Why It Fits
#1	Data Scientist Intern (AI Focus)	Exact match — production ML, AI tools, crypto domain
#2	ML Engineer Intern	End-to-end pipeline, MLflow + Docker stack
#3	Junior Data Scientist	Production deployments set you above average new grad
#4	Applied Data Science Intern	WorldQuant + Codveda + portfolio backs this fully
#5	Quantitative Data Analyst — Crypto	Crypto AI project + time-series is the differentiator
#6	AI Research Intern	Neural net (97.97% MNIST) + NLP pipeline
#7	NLP Engineer Intern	5-class classification, TF-IDF, macro-F1 evaluation
#8	ML Pipeline Engineer (Junior)	Docker + MLflow + CI/CD + Railway deployment
#9	Data Intelligence Analyst — Web3	Crypto domain + forecasting + pipeline overlap
#10	AI/Data Product Analyst	Bridges model work with market intelligence reporting

1. Production-deployed ML systems (rare at intern level)	Two live deployments — Crypto AI Intelligence System (Railway + Docker + CI/CD) and the Food Price Forecasting Streamlit app. Most intern candidates have Kaggle notebooks. You have shipped systems.
2. Full-stack ML ownership	From 50,000+ raw rows across 6 data sources through feature engineering, multi-model evaluation (SARIMA, XGBoost, Prophet), and Streamlit deployment — end-to-end, not just the modelling step.
3. Crypto / Web3 domain specificity	Your Crypto AI Intelligence System proves genuine domain engagement. For a Web3-focused role, this is the single most important signal — it is not a course project, it is a live system.
4. Quantified model evaluation discipline	MAPE 14.7% (best), 30.8% average; MNIST 97.97% vs. 92.49% baseline (+5.48pp); macro-F1 across 5 NLP classes. You report performance honestly and with context.
5. Native AI-augmented workflow	Claude, GitHub Copilot, and Jupyter AI are in your skills section and your daily practice. The Zoo Finance JD explicitly calls this out as desirable.
6. Remote work track record	Codveda internship — fully remote, 6 ML projects delivered asynchronously. Removes a common concern about early-career remote candidates.
7. Technical leadership under real constraints	Led a 6-person capstone team: pipeline coordination, conflict management, production deadline. That is leadership, not just a group project.
8. Technical breadth with demonstrated depth	Python, SQL, TensorFlow, NLTK, MLflow, Docker, Streamlit, CI/CD — every item traceable to a deployed system or quantified project outcome.

Current Strengths

- Python (strong) — production systems, ML pipeline, EDA, NLP
- Time-series modeling — SARIMA, SARIMAX, Prophet, XGBoost with quantified MAPE
- Deep learning basics — TensorFlow/Keras, multilayer dense nets, 97.97% MNIST accuracy
- NLP pipeline — TF-IDF, 5-class classification, macro-F1, stratified fold validation
- MLOps fundamentals — MLflow, Docker, Railway, CI/CD, Streamlit
- AI tool integration — Claude, GitHub Copilot, Jupyter AI in daily workflow
- End-to-end project ownership and production deployment
- SQL — functional (joins, aggregations); not yet advanced

Identified Skill Gaps (Web3 / DeFi Focus)

Gap	Severity	Note
On-chain analytics tools (Dune Analytics, Flipside, Nansen)	LOW	Explicitly listed in JD — start this week
Web3/DeFi fundamentals (smart contracts, liquidity, protocols)	MEDIUM	Needed to ask the right analytical questions
Advanced SQL (window functions, CTEs, crypto datasets)	MEDIUM	Functional now; crypto queries get complex
LLM/RAG integration	MEDIUM	Increasingly expected in AI-focus roles
Excel/Google Sheets proficiency	MEDIUM	Listed in JD — a quick win, close in one afternoon
Statistical depth (Bayesian inference, A/B testing)	MEDIUM	WorldQuant Lab helps; needs active focus
On-chain/Web3 API integration (CoinGecko, Etherscan)	MEDIUM	Your crypto project is the natural place to close this
Advanced data visualization (Plotly Dash, D3.js)	LOW	Streamlit dashboards functional; deeper viz would help

90-Day Learning Roadmap

NOW — 4 Weeks

- Dune Analytics (free) — write 3 original on-chain SQL queries on a DeFi protocol; put them on GitHub
- Alchemy University 'Web3 Fundamentals' — free, ~4 hours; closes vocabulary gap immediately
- Excel pivot tables + VLOOKUP — spend 2 hours, close the JD gap in an afternoon
- Integrate CoinGecko or Etherscan API into Crypto AI System Phase 3

1–3 Months

- Advanced SQL: Mode Analytics free tutorial — window functions and CTEs, practiced on crypto datasets
- LLM basics: Hugging Face NLP Course fine-tuning modules (you already have the TF-IDF foundation)

- Bayesian inference module (probabilistic ML) — fills statistical depth gap
- On-chain API integration live in deployed Crypto AI System

3–6 Months

- Build a Web3 market intelligence dashboard: Dune + Python + Streamlit/Dash — flagship portfolio piece
- Contribute to an open-source DeFi analytics project on GitHub — signals genuine community engagement
- Advanced visualization — Plotly Dash or D3.js
- Build one production LLM/RAG application — chain a language model to real data

Unique Value Proposition

You ship. At an experience level where most candidates are still doing guided labs, you have two production ML systems with real architecture decisions. You work AI-natively in a space (data science + Web3) where that multiplies output. Your value is not potential — it is demonstrated delivery at intern level with junior-level architecture decisions.

Professional Positioning Statement

"A data scientist who doesn't just model — she builds and deploys. From raw multi-source data engineering through multi-model evaluation to live production deployments, Deborah's portfolio lives in the real world: a Crypto AI Intelligence System running in production, a food price forecasting system serving Northeast Nigeria, and six ML projects delivered remotely for a tech internship. Currently deepening ML engineering and Web3 analytics expertise through AWS AI & ML Scholars, DataCamp ML Engineering, and WorldQuant's Applied Data Science Lab."

Recruiter-Ready LinkedIn Headline Options

Option A

Data Scientist & ML Engineer | Production Crypto AI + Forecasting Systems | Python · MLflow · Docker · Streamlit | CS Graduate | Remote

Option B

Building ML Systems That Deploy | Crypto AI · Time-Series Forecasting · NLP | CS Graduate | Open to Remote Data Science Roles

Option C

Data Scientist | Crypto AI Systems · DeFi Data Analytics · Production ML | Python · Docker · AI-Augmented Workflow | CS Graduate

Three Most Impressive Achievements to Lead With

1. Crypto AI Intelligence System

Live deployment, multi-source ingestion (Phases 1-2), Phase 3 ML pipeline in active development. Demonstrates initiative, ML engineering depth, and crypto domain knowledge that is not faked.

2. Early Warning Food Price Monitor

50,000+ rows to production. Team lead role. 14.7% best MAPE across 13 states and 7 commodities. Multi-source integration (WFP, CBN, NBS, ACLED, FAO) and multi-model head-to-head evaluation.

3. MNIST Neural Network — 97.97% Accuracy

5.48 percentage-point improvement over baseline (92.49%). Architecture decisions: ReLU, dropout, batch normalization, softmax. Shows model design understanding, not just model fitting.

Behavioral Stories Mapped to Common Questions

Question	Story / Anchor
Tell me about messy real-world data	Food Price Monitor — 50K rows, 6 inconsistent sources, multi-year alignment
Describe leading a technical team	Capstone lead — pipeline coordination, conflict management, production deadline
How do you use AI tools in your workflow?	Claude + GitHub Copilot + Jupyter AI daily — with specific examples of where AI a
Walk me through a project to deployment	Crypto AI System — architecture decisions, tool selection, what Phases 1-2 includ
Why Web3 / DeFi specifically?	The Crypto AI System is your evidence. You were building in this space before yo

Probe Areas — Prepare Specific Answers

Q: 'What does production deployed actually mean for the Crypto AI System?'

A: Be specific: Railway hosting, Docker container, CI/CD pipeline, Streamlit dashboard contents, which data sources are live vs. in development. Do not overstate Phase 3.

Q: 'Have you used Dune Analytics, Flipside, or Nansen?'

A: If not yet: 'I haven't used them extensively but I've been actively exploring Dune since seeing it in the JD — I ran [X] queries this week.' This is why you should start Dune this week regardless.

Q: 'You listed SQL — how advanced is it?'

A: Be honest: 'Functional — joins, aggregations, basic window functions. I'd describe it as intermediate-basic, and it's something I'm actively deepening on crypto datasets.'

Q: 'Why Web3/DeFi — tell me your on-chain knowledge'

A: Use your Crypto AI System as proof of engagement, then be transparent about on-chain tooling depth and the steps you are actively taking to build it.



5 STRATEGIC TARGET COMPANIES

Selected for: direct profile alignment | realistic entry at your level | growth trajectory | compensation | remote accessibility from Lagos, Nigeria

Coinbase

Directly aligned with your Crypto AI Intelligence System. Coinbase has a serious data science function analyzing on-chain behavior, trading patterns, and user behavior at scale. Remote-friendly, competitive pay (\$40-60/hr intern range), and your live crypto project is a direct talking point.

Your angle: Live crypto data pipeline + ML system + DeFi economics interest.

Immediate action: Follow Coinbase Engineering blog. Build an analysis of a Coinbase asset using your crypto system. Apply to their intern/new grad data pipeline.

Hugging Face

ML-first, remote-first company that values practical builders over paper credentials. Strong track record of hiring early-career engineers who contribute to the ecosystem first. Your NLP experience and AI tool usage puts you in their orbit.

Your angle: NLP background + Python ML engineering + active AI tool user.

Immediate action: Find a Hugging Face dataset or model card overlapping your work; submit a PR or model card improvement. Then apply formally.

Chainalysis

Leading blockchain data and intelligence company — their entire product is on-chain data analysis, exactly what your Crypto AI System points toward. Growing fast, well-funded, strong compensation, remote roles available.

Your angle: Crypto AI System + data pipeline engineering + on-chain market intelligence interest.

Immediate action: Study their public DeFi market reports; reference findings in your cover letter. Start Dune Analytics this week — it is how their analysts work.

Scale AI

Works at the intersection of data quality and AI model development. Roles for early-career ML engineers focused on data pipelines, evaluation, and model performance. Strong pay for early career, remote options, directly adjacent to your MLOps stack.

Your angle: Production ML pipelines + model evaluation discipline (MAPE/F1 with baselines) + MLflow/Docker/CI/CD.

Immediate action: Apply to their AI trainer or data quality engineer programs as an entry point — these feed into core ML pipeline teams.

Andela

Places African tech talent into senior remote roles at top global companies. Getting accepted means your profile is reviewed by global hiring managers who have already committed to remote-inclusive hiring from Africa. For a Lagos-based candidate, this is a force multiplier — one application, multiple doors.

Your angle: CS graduate + production ML systems + remote work experience + growing ML engineering skills.

Immediate action: Apply to Andela's talent network directly. Polish GitHub and LinkedIn before doing so — they evaluate portfolio quality carefully.

Strategic Note on Microsoft, Google, Amazon & Meta

These are achievable medium-term targets — not immediate reaches — once you have 6-12 more months of production portfolio depth and ideally one of the 5 companies above on your resume. Microsoft (LEAP program), Google (STEP), Amazon (ML internships), Meta (production engineering tracks) are all realistic progression targets. Keep building — the portfolio you have now is the foundation.

Deborah Olofin | The Midnight Architect

github.com/Mi-kami | linkedin.com/in/deborah-olofin | deboraholofin43@gmail.com

60 Days Claude Challenge — Day 12 of 60 | June 2026